Context-Dependent Language Understanding

Alane Suhr Cornell University

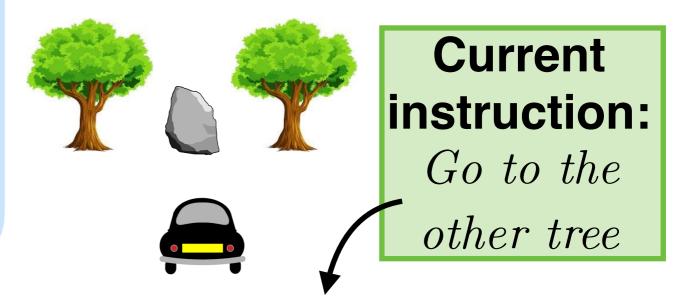
Context:

Previous instruction:

Go to the tree on the right

Previous interpretation:

 $\lambda a.move(a) \wedge \lambda x.tree(x) \wedge right-$ of(x, rock) \wedge to(a, x)



 $\lambda a.move(a) \land \lambda x.tree(x) \land \neg right-of(x, rock) \land to(a, x)$

Context:

Previous instructions

Previous instruction:

Go to the tree on the right

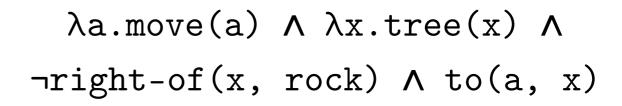
Previous interpretation:

 $\lambda a.move(a) \wedge \lambda x.tree(x) \wedge right-$ of(x, rock) \wedge to(a, x)



Current instruction:

Go to the other tree



Context:

- Previous instructions
- Previous interpretations of instructions

Previous instruction:

Go to the tree on the right

Previous interpretation:

 $\lambda a.move(a) \wedge \lambda x.tree(x) \wedge right-$ of(x, rock) \wedge to(a, x)



Current instruction:

Go to the other tree

 $\lambda a.move(a) \land \lambda x.tree(x) \land \neg right-of(x, rock) \land to(a, x)$

Context:

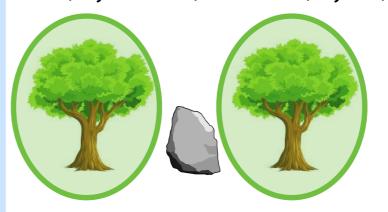
- Previous instructions
- Previous interpretations of instructions
- Current world state

Previous instruction:

Go to the tree on the right

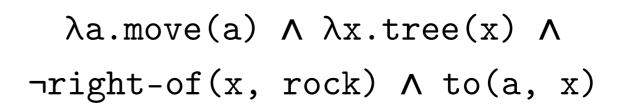
Previous interpretation:

 $\lambda a.move(a) \wedge \lambda x.tree(x) \wedge right-$ of(x, rock) \wedge to(a, x)



Current instruction:

Go to the other tree



Context:

- Previous instructions
- Previous interpretations of instructions
- Current world state

Previous instruction:

Go to the tree on the right

Previous interpretation:

 $\lambda a.move(a) \wedge \lambda x.tree(x) \wedge right-$ of(x, rock) \wedge to(a, x)



 $\lambda a.move(a) \land \lambda x.tree(x) \land \neg right-of(x, rock) \land to(a, x)$

Context:

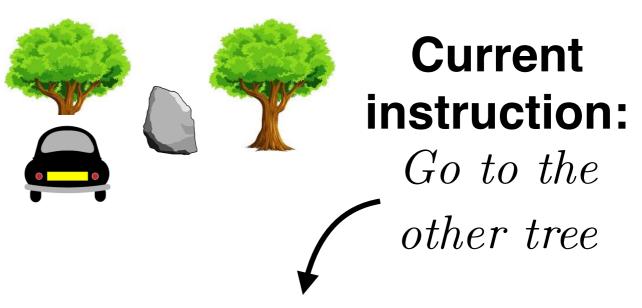
- Previous instructions
- Previous interpretations of instructions
- Current world state
- Previous actions taken by user or agent
- Previous world states

Previous instruction:

Go to the tree on the right

Previous interpretation:

 $\lambda a.move(a) \wedge \lambda x.tree(x) \wedge right-$ of(x, rock) \wedge to(a, x)



 $\lambda a.move(a) \land \lambda x.tree(x) \land \neg right-of(x, rock) \land to(a, x)$

Outline

- Case study: natural language interface to a database
 - A. Previous utterances
 - B. Previous interpretations of utterances
- 2. Situated interaction context
- 3. Open questions

User | Show me flights from Seattle to Boston next Monday

User Show me flights from Seattle to Boston next Monday

User Show me flights from Seattle to Boston next Monday

```
SQL (SELECT DISTINCT flight_flight_id FROM flight WHERE
        (flight.from_airport IN (SELECT airport_service.airport_code
Query FROM airport_service WHERE airport_service.city_code IN
        (SELECT city_code FROM city WHERE city_city_name =
        'SEATTLE'))) AND (flight.to_airport IN (SELECT
        airport_service.airport_code FROM airport_service WHERE
        airport_service.city_code IN (SELECT city.city_code FROM
        city WHERE city.city_name = 'BOSTON'))) AND
        (flight_flight_days IN (SELECT days.days_code FROM days
        WHERE days.day_name IN (SELECT date_day.day_name FROM
        date_day WHERE date_day.year = 1993 AND
        date_day.month_number = 2 AND date_day.day_number = 8)));
```

User | Show me flights from Seattle to Boston next Monday

User On American Airlines

[Hemphill et al. 1990, Dahl et al. 1994]

User Show me flights from Seattle to Boston next Monday

Result Found 31 Flights: ARAMA

User On American Airlines

Result Found 5 Flights: ARAMA

User Show me flights from Seattle to Boston next Monday

Result
User On American Airlines

Result
User Which ones arrive after 7pm?

Result
No flights found.

User | Show me flights from Seattle to Boston next Monday Result Found 31 Flights: 77777 User On American Airlines Result Found 5 Flights: 7777 User Which ones arrive after 7pm? Result No flights found. User | Show me Delta flights Result Found 5 Flights: 777

[Hemphill et al. 1990, Dahl et al. 1994]

ATIS

- Flight information, 27 tables, 162K entries
- Small corpus: <2000 interactions
- Complex queries: average 102.9 tokens each; 93% reference >3 tables
- Long interactions: average 7 turns; maximum: 64

Previous Approaches

- Model: statistical decision tree
- Representation: semantic frames
- Two-step process:
 - Map request to context-independent frame
 - 2. Use context to decide how to use past constraints
- Context: previous decisions on using constraints

Previous Approaches

- Model: weighted linear combinatory categorial grammar
- Representation: lambda calculus
- Two-step process:
 - Map request to context-independent expression
 - 2. Use context to map to context-dependent expression
- Context: previously generated expressions

Suhr et al. 2018

- As an interaction progresses, the meaning of an utterance becomes highly dependent on the meaning of previous utterances
- Interaction history: previously generated queries and previous requests

Challenges

- How do we make use of previous requests?
- How do we make use of previously generated queries?

Context from Previous Requests

Show me all flights from Boston to Pittsburgh on Wednesday of next week which depart from Boston after 5pm

: (3 turns)

User | Please describe the class of service Y

• (5 turns)

User | Show the cost of tickets on flight US 345

Context from Previous Requests

- Relevant but elided information mentioned multiple turns before
- User may change focus during interaction
- Solution: implicit mechanism for carrying information from beginning to end of interaction

Context from Previous Queries

User Show me flights from Seattle to Boston next Monday

```
SQL (SELECT DISTINCT flight.flight_id FROM flight WHERE
        (flight.from_airport IN (SELECT airport_service.airport_code
Query FROM airport_service WHERE airport_service.city_code IN
        (SELECT city_code FROM city WHERE city_city_name =
        'SEATTLE'))) AND (flight.to_airport IN (SELECT
        airport_service.airport_code FROM airport_service WHERE
        airport_service.city_code IN (SELECT city.city_code FROM
        city WHERE city.city_name = 'BOSTON'))) AND
        (flight_flight_days IN (SELECT days.days_code FROM days
        WHERE days.day_name IN (SELECT date_day.day_name FROM
        date_day WHERE date_day.year = 1993 AND
        date_day.month_number = 2 AND date_day.day_number = 8)));
```

Context from Previous Queries

User On American Airlines

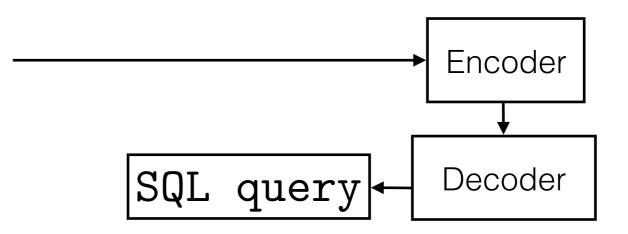
```
(SELECT DISTINCT flight.flight_id FROM flight WHERE
 SQL (flight.airline_code = 'AA') AND (flight.from_airport IN
Query (SELECT airport_service.airport_code FROM airport_service
        WHERE airport_service.city_code IN (SELECT city.city_code
        FROM city WHERE city.city_name = 'SEATTLE'))) AND
        (flight.to_airport IN (SELECT airport_service.airport_code
        FROM airport_service WHERE airport_service.city_code IN
        (SELECT city_code FROM city WHERE city_city_name =
        'BOSTON'))) AND (flight.flight_days IN (SELECT days.days_code
        FROM days WHERE days.day_name IN (SELECT date_day.day_name
        FROM date_day WHERE date_day.year = 1993 AND
        date_day.month_number = 2 AND date_day.day_number = 8))));
```

Context from Previous Queries

- Segments corresponding to earlier mentioned constraints and entities appear in later queries
- Solution: explicit mechanism for composing later
 SQL queries from segments of previous ones

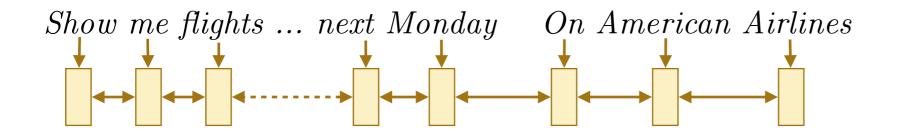
Semantic Parsing as Sequence-to-Sequence

Show me flights from Seattle to Boston next Monday

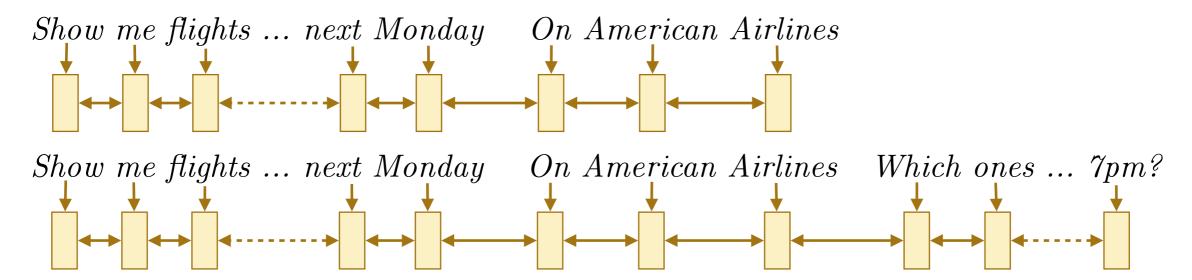


Semantic Parsing as Sequence-to-Sequence

Show me flights from Seattle to Boston next Encoder MondaySQL query Decoder On American Airlines Encoder Decoder SQL query



Simple approach: encode interaction history with a bidirectional RNN



Simple approach: encode interaction history with a bidirectional RNN

Downside: need to re-run bidirectional encoder with each new utterance

Downside: need to re-run bidirectional encoder with each new utterance

Imposes tradeoff between efficiency and access to interaction history

Downside: need to re-run bidirectional encoder with each new utterance

Imposes tradeoff between efficiency and access to interaction history

n = 0

Previous constraints not available to decoder

n, number of previous concatenated utterances

Downside: need to re-run bidirectional encoder with each new utterance

Imposes tradeoff between efficiency and access to interaction history

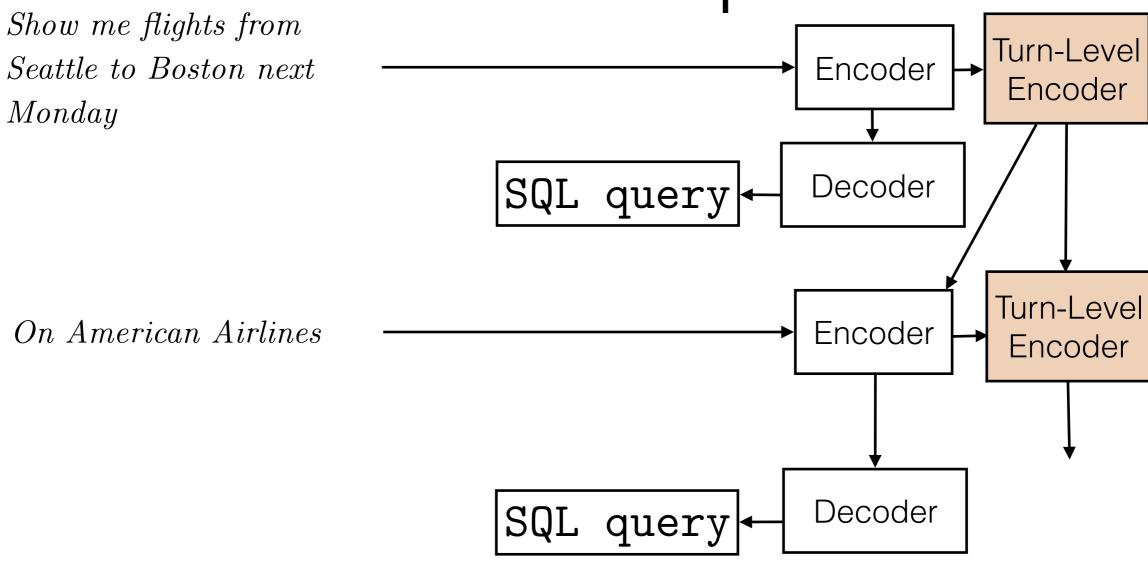
$$n = 0$$

Previous constraints not available to decoder

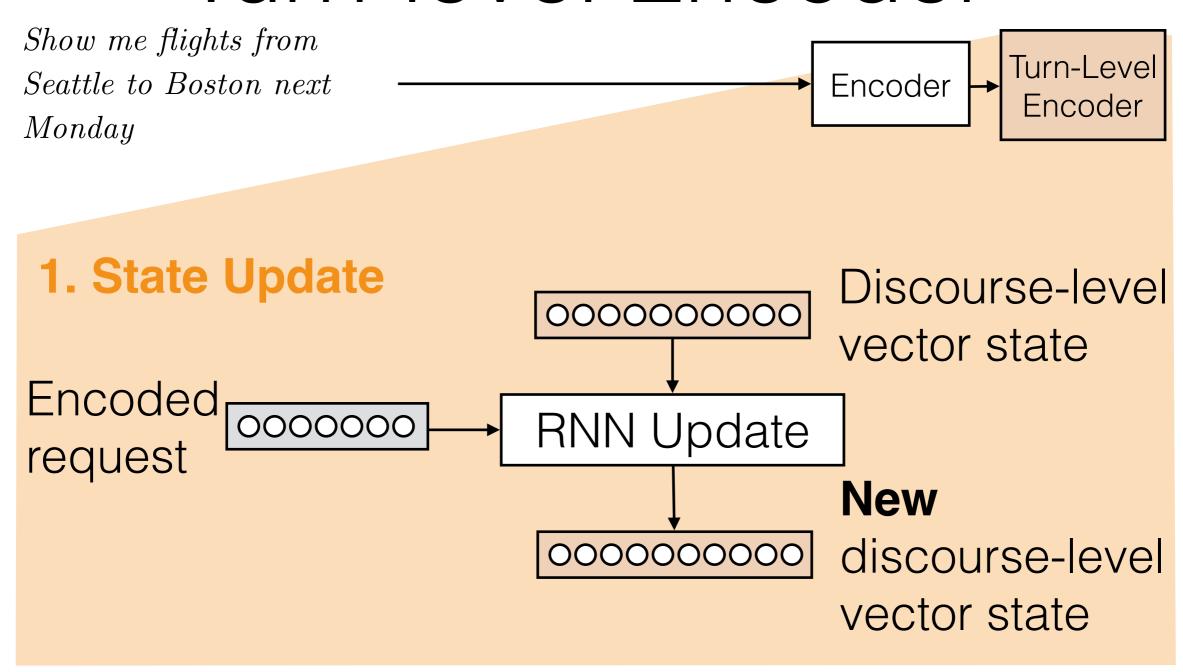
$$n = \infty$$

Inefficient to encode — must re-encode all utterances

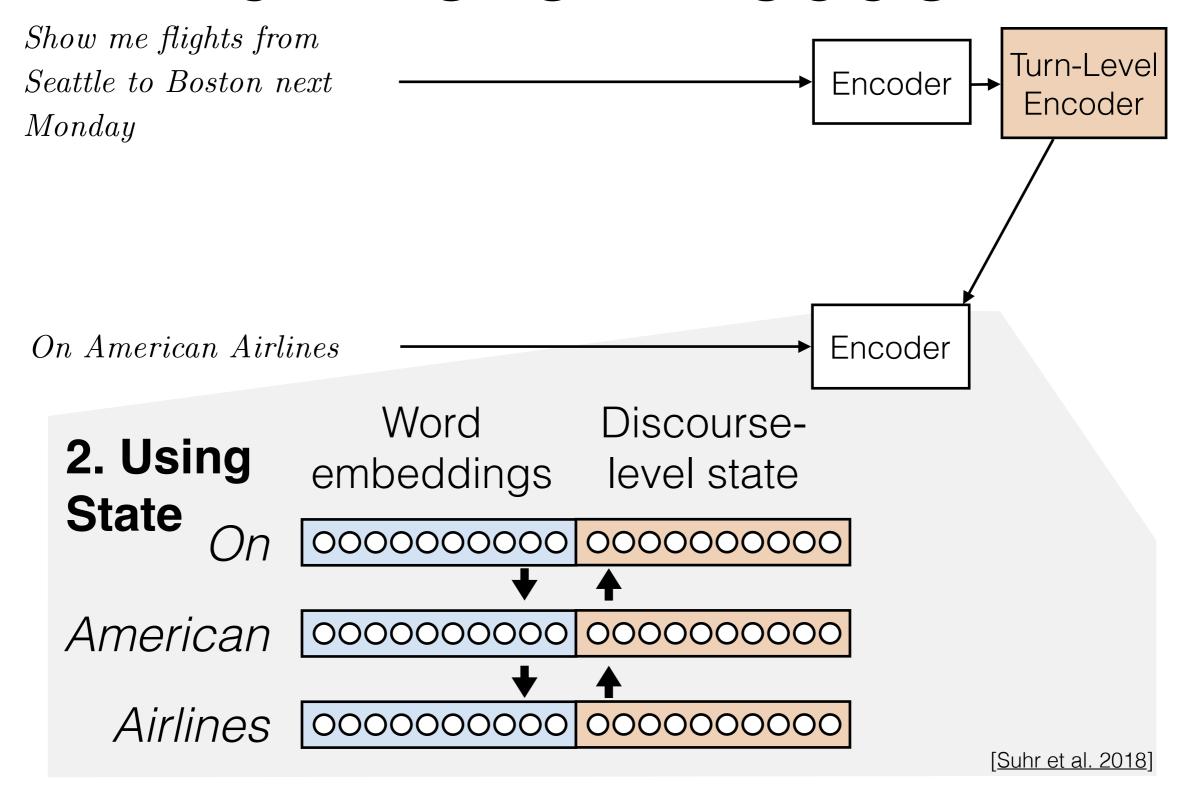
n, number of previous concatenated utterances



Turn-level Encoder



Turn-level Encoder



Turn-level Encoder

How can we make use of previous requests?

- Efficiently? —> Only run bidirectional RNN over newest utterance
- Using all information from interaction history?
 - —> Vector state persistent through interaction

Challenges

Challenges

- How do we make use of previous requests?
- How do we make use of previously generated queries?

Turn-level encoder

Incorporating Previous Queries

Show me flights from Seattle to Boston next Encoder MondaySQL query Decoder Encoder On American Airlines **Query Segment** Copying Decoder SQL query

Query Segment Copying

Previous Query:

(SELECT DISTINCT flight.flight_id FROM flight
WHERE (flight.from_airport IN (SELECT
airport_service.airport_code FROM airport_service
WHERE airport_service.city_code IN (SELECT
city.city_code FROM city WHERE city.city_name =

Decoder

1. Segment Extraction

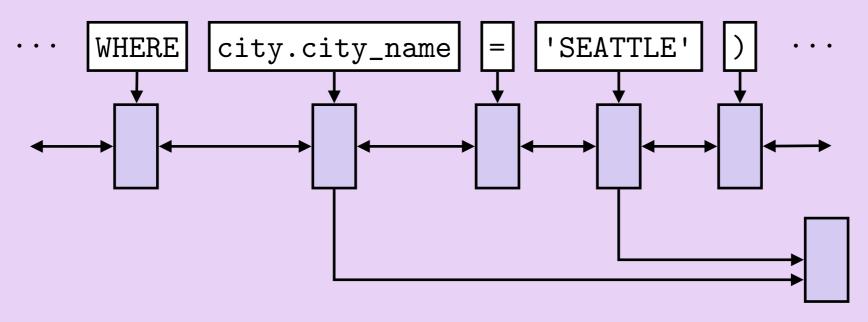
```
city.city_name = 'SEATTLE'
city.city_name = 'BOSTON'
date_day.year = 1993
date_day.month_number = 2
date_day.day_number = 8
```

Deterministic, operates on the SQL tree

Query Segment Copying

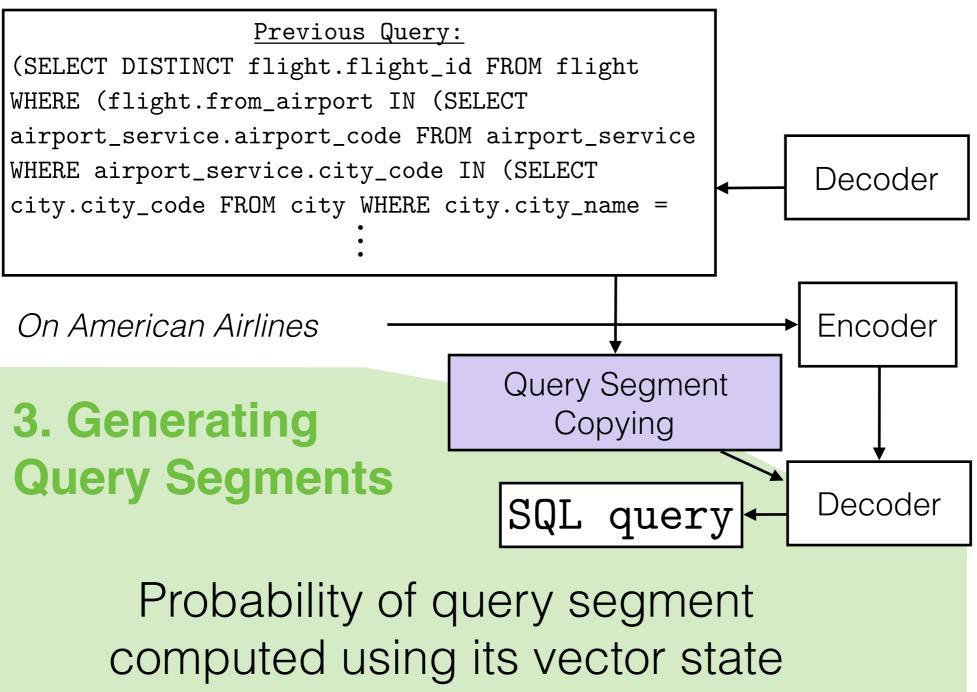
Previous Query: (SELECT DISTINCT flight.flight_id FROM flight WHERE (flight.from_airport IN (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN (SELECT city.city_code FROM city WHERE city.city_name =

2. Segment Encoding



city.city_name = 'SEATTLE'

Query Segment Copying



Challenges

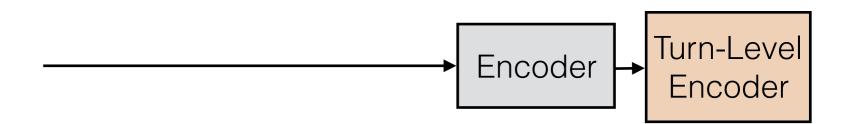
Challenges

How do we make use of previous requests?

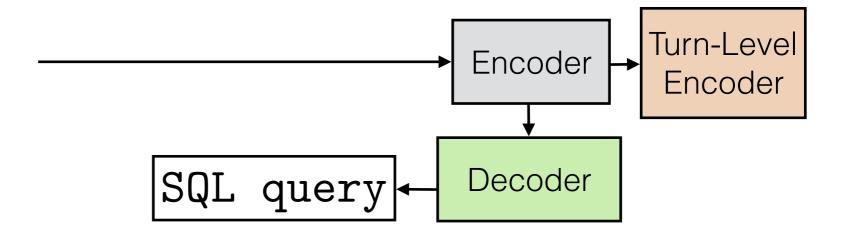
 How do we make use of previously generated queries? Turn-level encoder

Query segment copying

Show me flights from Seattle to Boston next Monday



Show me flights from Seattle to Boston next Monday



Show me flights from Seattle to Boston next Monday

SQL query Decoder

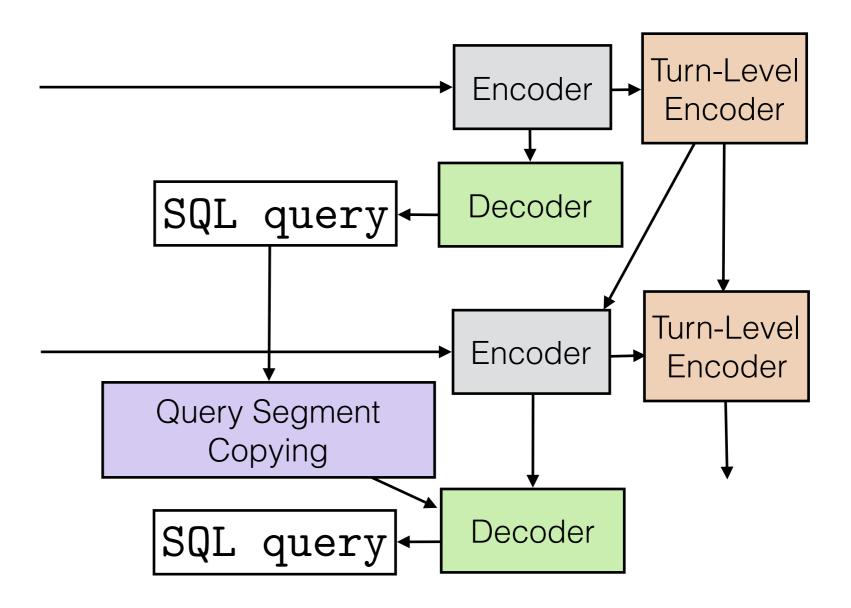
Encoder

Encoder

On American Airlines

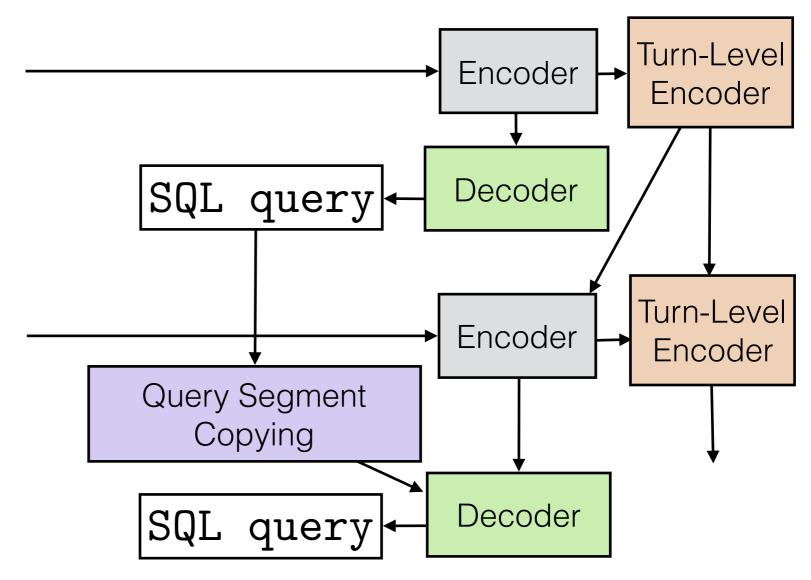
Show me flights from Seattle to Boston next Monday

On American Airlines



Show me flights from Seattle to Boston next Monday

On American Airlines



- Gradually build computation graph for entire interaction
- Directly predict context-dependent representations

Learning

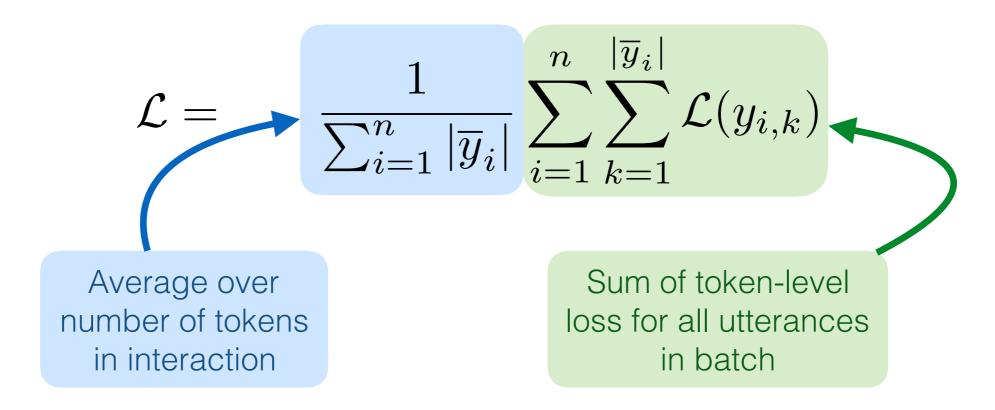
- Training data: interactions with request-SQL pairs
- Objective: minimize token-level cross-entropy loss

Interaction-Level Learning

- Each batch consists of an entire interaction
- Build one computation graph, and backpropagate through entire interaction
- Each batch contains a different number of output sequences
- This introduces several learning challenges

Interaction Loss

Loss for an interaction y



Problem: sequences in longer interactions have relatively lower gradient due to token-level averaging

Batch Re-Weighting

Loss for an interaction y with batch re-weighting

$$\mathcal{L} = \frac{n}{B} \frac{1}{\sum_{i=1}^{n} |\overline{y}_i|} \sum_{i=1}^{n} \sum_{k=1}^{|\overline{y}_i|} \mathcal{L}(y_{i,k})$$
 Re-normalized based on length of interaction

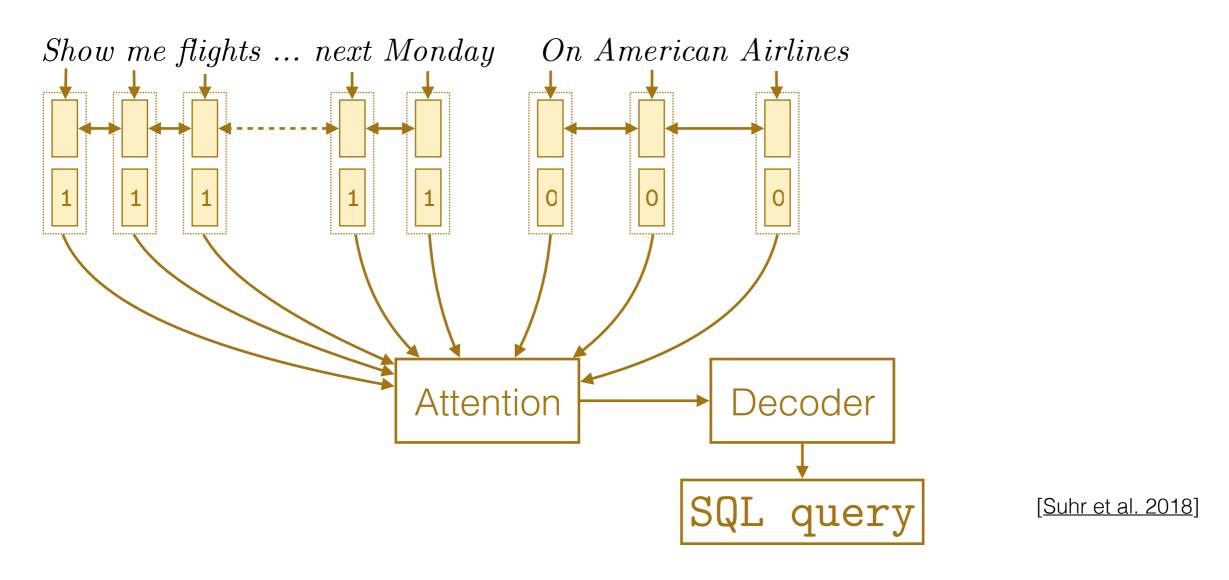
Solution: gradient will have similar magnitude regardless of interaction length

Attention on Interaction History

- Attention is effective for sequence-to-sequence models
- Our model: attend over last few utterances
- Problem: hidden states don't encode which utterance they are coming from

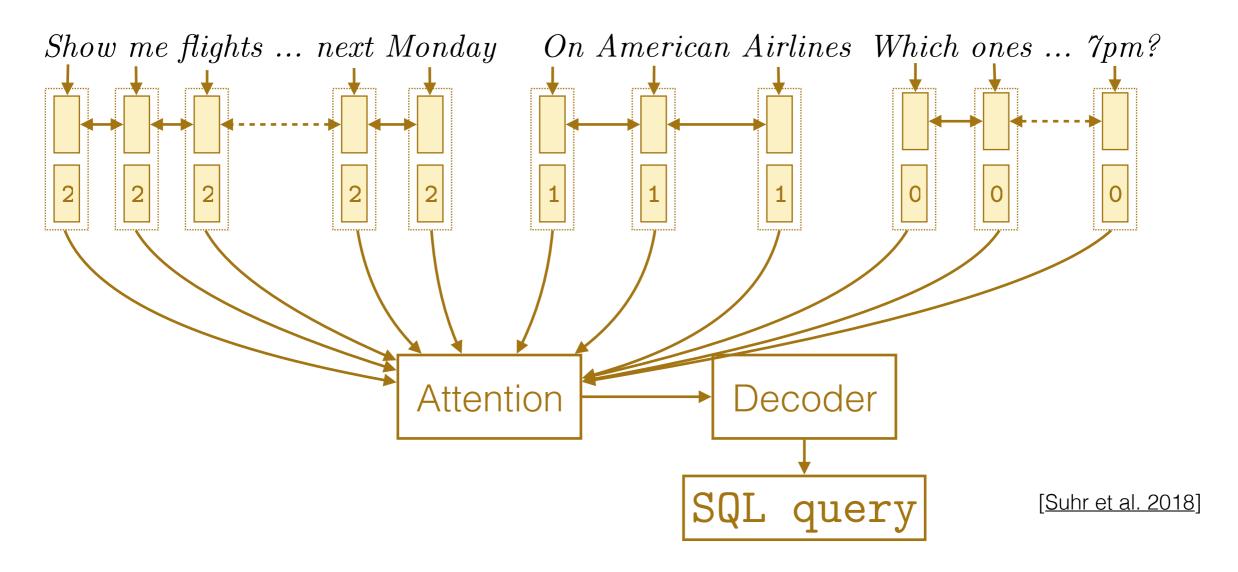
Attention on Interaction History

Solution: add positional embeddings to input hidden states in attention computations



Attention on Interaction History

Solution: add positional embeddings to input hidden states in attention computations



Learning to Copy Segments

- Goal: model should learn to copy as much as possible, when appropriate
- Challenge: tokens are not annotated with whether they are generated or copied
- Solution: heuristic extraction and alignment of segments

Previous Query:

(SELECT DISTINCT flight.flight_id
FROM flight WHERE
(flight.from_airport IN (SELECT
airport_service.airport_code FROM
airport_service WHERE
airport_service.city_code IN (SELECT
city.city_code FROM city WHERE
city.city_name =

```
city.city_name = 'SEATTLE'
city.city_name = 'BOSTON'
date_day.year = 1993
date_day.month_number = 2
date_day.day_number = 8
:
```

Includes all constituents of the SQL tree

<u>Request</u>

On American Airlines

Gold Query

```
(SELECT DISTINCT flight.flight_id
FROM flight ...
city.city_name = 'SEATTLE'))) ...
```

Available Segments

```
1. city.city_name = 'SEATTLE'
```

2. city.city_name = 'BOSTON'

 $3. date_day.year = 1993$

4. date_day.month_number = 2

5. date_day.day_number = 8
:

Request On American Airlines

```
Gold Query
(SELECT DISTINCT flight.flight_id
FROM flight ...
city.city_name = 'SEATTLE'))) ...
```

```
Available Segments

1. city.city_name = 'SEATTLE'

2. city.city_name = 'BOSTON'

3. date_day.year = 1993

4. date_day.month_number = 2

5. date_day.day_number = 8

:
```

```
Request
On American Airlines
```

```
Gold Query

(SELECT DISTINCT flight.flight_id

FROM flight ...

SEGMENT#1

))) ...
```

```
Available Segments

1. city.city_name = 'SEATTLE'

2. city.city_name = 'BOSTON'

3. date_day.year = 1993

4. date_day.month_number = 2

5. date_day.day_number = 8

:
```

Replace identified segment with reference

<u>Request</u>

From Seattle to Denver

Gold Query

```
(SELECT DISTINCT flight.flight_id
FROM flight ...
city.city_name = 'SEATTLE')) ...
```

```
Available Segments
```

```
1. city.city_name = 'SEATTLE'
```

2. city.city_name = 'BOSTON'

 $3. date_day.year = 1993$

4. date_day.month_number = 2

5. date_day.day_number = 8
:

```
Request From Seattle to Denver
```

```
Gold Query
(SELECT DISTINCT flight.flight_id
FROM flight ...
city.city_name = 'SEATTLE'))) ...
```

```
Available Segments

1. city.city_name = 'SEATTLE'

2. city.city_name = 'BOSTON'

3. date_day.year = 1993

4. date_day.month_number = 2

5. date_day.day_number = 8

:
```

Don't replace: entity present in current request

Request

From Seattle to Denver

Gold Query

```
(SELECT DISTINCT flight.flight_id
FROM flight ...
city.city_name = 'SEATTLE')) ...
```

Available Segments

```
1. city.city_name = 'SEATTLE'
```

2. city.city_name = 'BOSTON'

 $3. date_day.year = 1993$

4. date_day.month_number = 2

5. date_day.day_number = 8
:

Don't replace: entity present in current request

Experiments

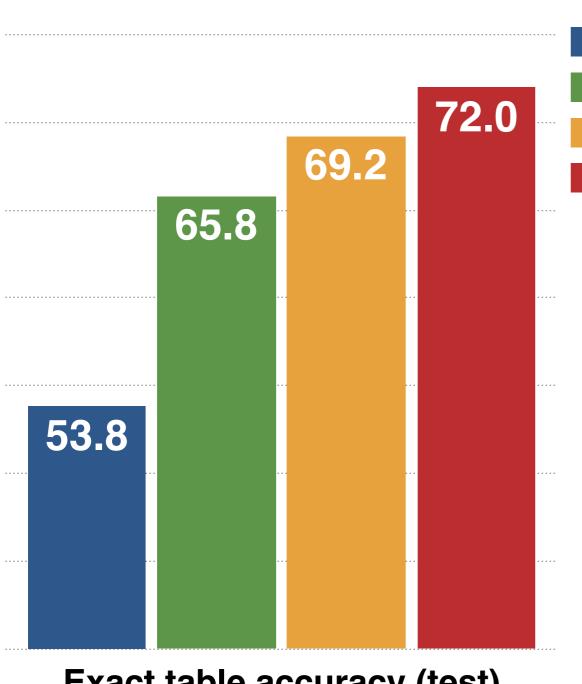
- Seq2Seq w/o history seq2seq on current utterance only
- Seq2Seq + history seq2seq by concatenating last four utterances
- Full model
 use turn-level encoder
 and query segment copying

Evaluation metric: Exact table accuracy

 Measure effect of error propagation: full model with access to gold previous query

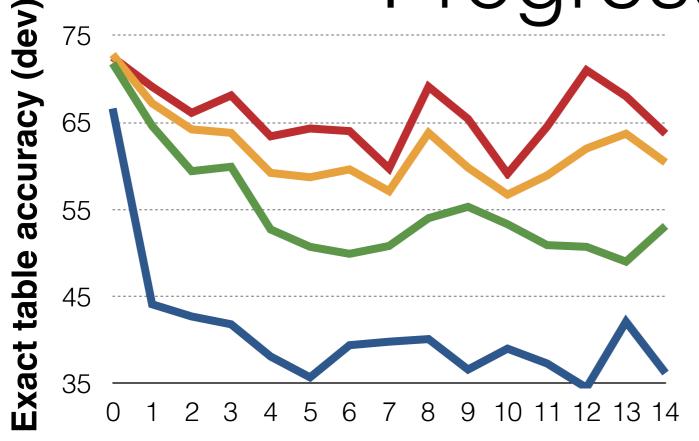
[Suhr et al. 2018]

Test Results



- Seq2Seq w/o history
- Seq2Seq + history
- Full model
- Full model (with gold previous query)
 - Using interaction history is critical
 - Error propagation contributes about 3% performance drop

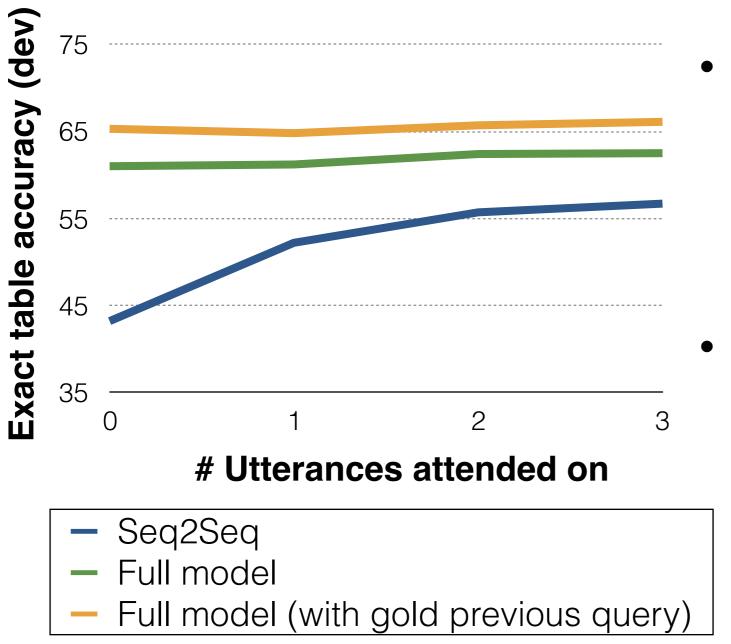
Performance as Interactions Progress



- Without interaction history, performance drops immediately
- Our model: relatively stable

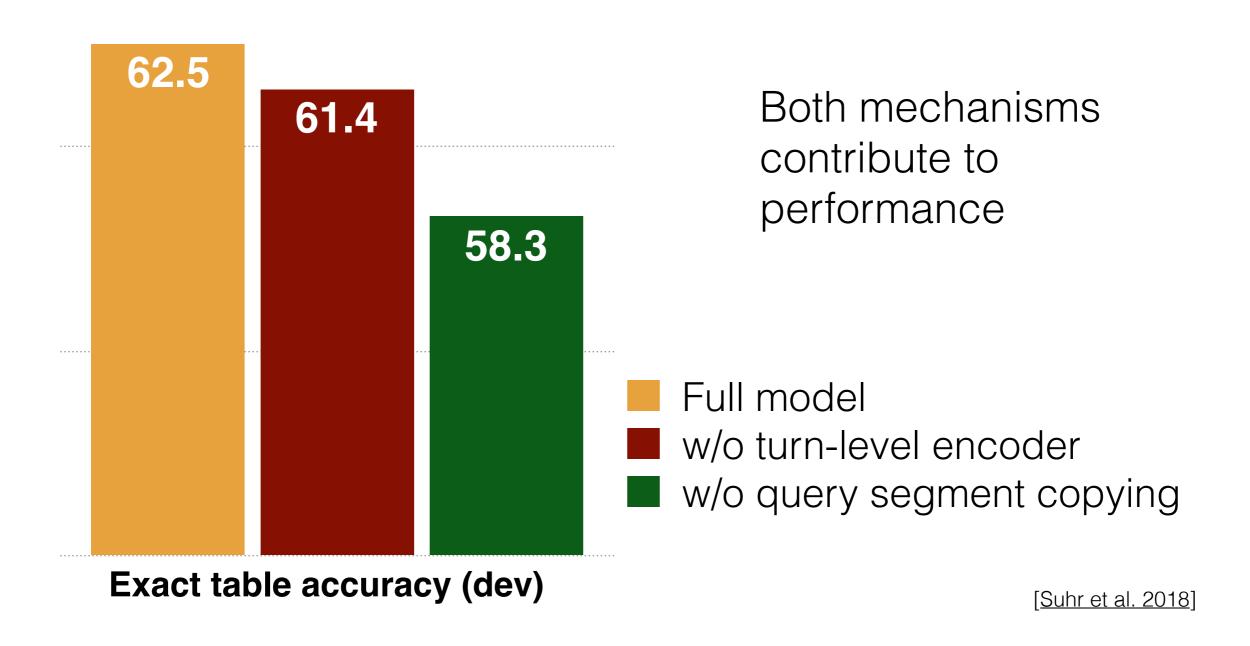
- Turn index in interaction
- Seq2Seq w/o history
- Seq2Seq + history
- Full model
- Full model (with gold previous query)

Attention Analysis

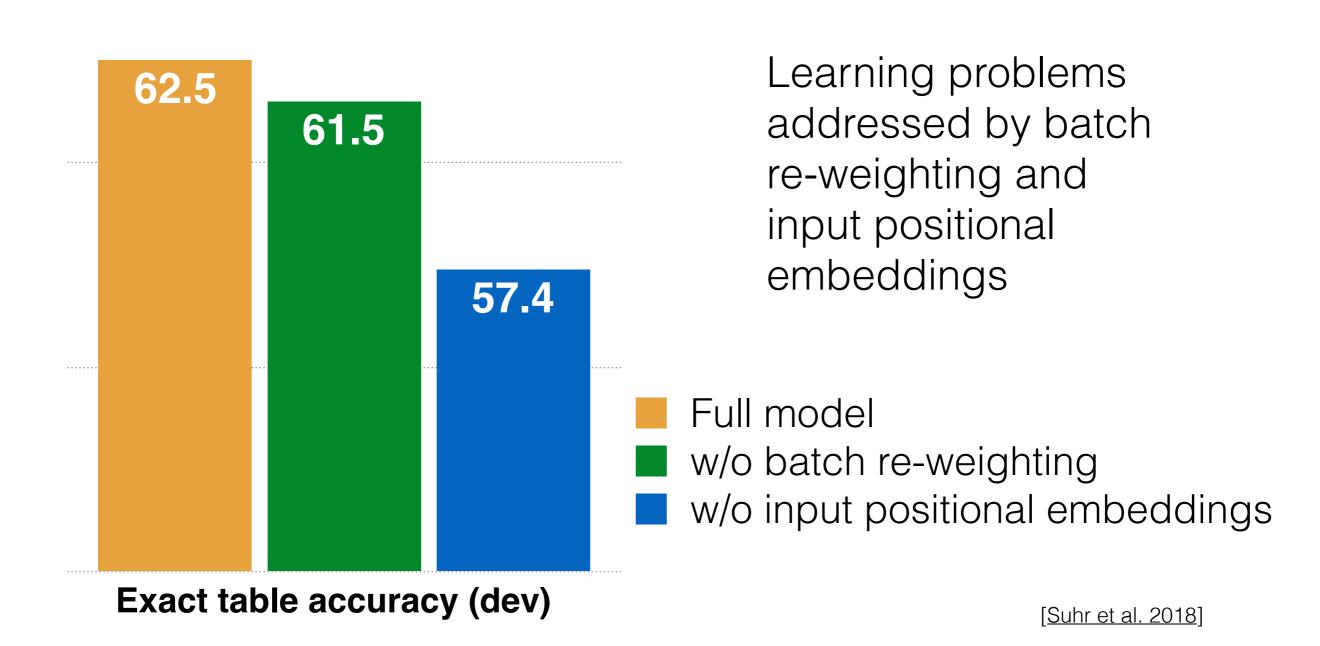


- With our mechanisms, can still perform well while only attending on current utterance
 - Seq2seq models need to attend over more utterances inefficient

Mechanism Ablations



Turn-Level Encoder Ablations



Error Propagation

User Which ones arrive around 7pm?

```
( SELECT DISTINCT flight.flight_id FROM flight WHERE
       ( flight.from_airport IN ( SELECT
Query airport_service.airport_code FROM airport_service WHERE
       airport_service.city_code IN ( SELECT city.city_code FROM
       (flight.to_airport IN (SELECT airport_service.airport_code
       FROM airport_service WHERE airport_service.city_code IN
        ( SELECT city.city_code FROM city WHERE city.city_name =
        'BALTIMORE' ) ) AND ( flight.flight_days IN ( SELECT
       days.days_code FROM days WHERE days.day_name IN ( SELECT
       date_day.day_name FROM date_day WHERE date_day.year = 1991
       AND date_day.month_number = 9 AND date_day.day_number =
       6 ) ) AND (flight.arrival_time >= 1630 AND
       flight.arrival_time <= 1730 ) ) ) );
                                                       [Suhr et al. 2018]
```

Error Propagation

User Which ones arrive around 7pm?

SQL Query

> Error: looking for flights around 5pm

flight.arrival_time >= 1630 AND

flight.arrival_time <= 1730

Error Propagation

User | Which kind of airplane is that?

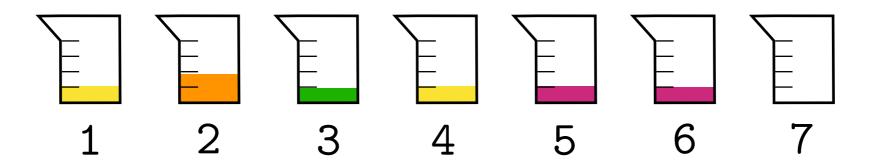
```
( SELECT DISTINCT aircraft.aircraft_code FROM aircraft WHERE
        aircraft.aircraft_code IN ( SELECT
Query equipment_sequence.aircraft_code FROM equipment_sequence
        WHERE equipment_sequence.aircraft_code_sequence IN ( SELECT
        flight.aircraft_code_sequence FROM flight WHERE
        ( flight.arrival_time >= 1630 AND flight.arrival_time <=</pre>
        1730 AND (flight.from_airport IN (SELECT
        airport_service.airport_code FROM airport_service WHERE
        airport_service.city_code IN ( SELECT city.city_code FROM
        city WHERE city.city_name = 'ATLANTA' ) )
```

Code and Data

- Code available: https://github.com/clic-lab/atis
- Preprocessed data available:
 - https://github.com/clic-lab/atis_data
 - Requires LDC93S5, LDC94S19, and LDC95S26
 - Email <u>suhr@cs.cornell.edu</u> for access

Outline

- Case study: natural language interface to a database
 - A. Previous instructions
 - B. Previous interpretations of instructions
- 2. Situated interaction context
 - 3. Open questions

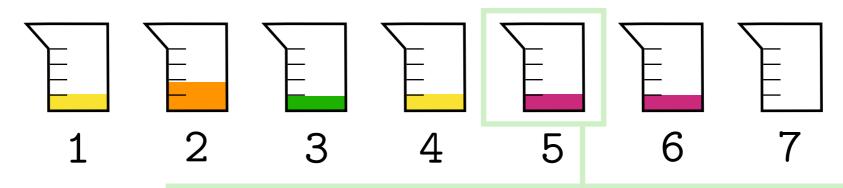


Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it

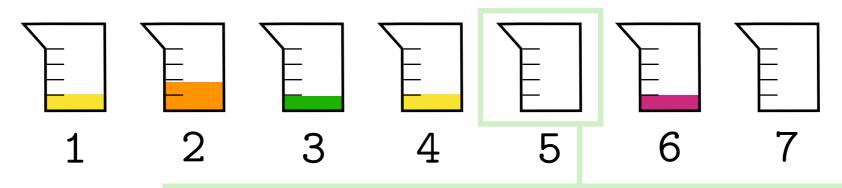


Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it

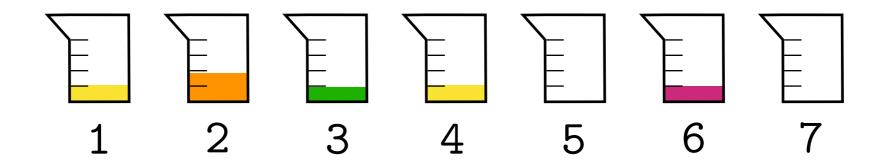


Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it

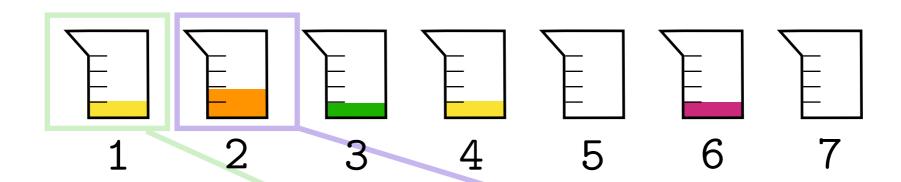


Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it



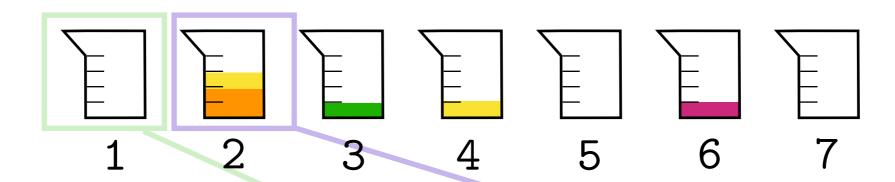
Empty out the leftmost beaker of purple chemical



Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it



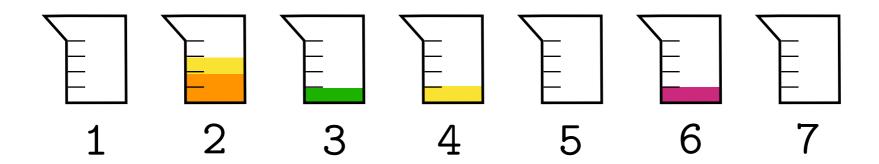
Empty out the leftmost beaker of purple chemical



Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it

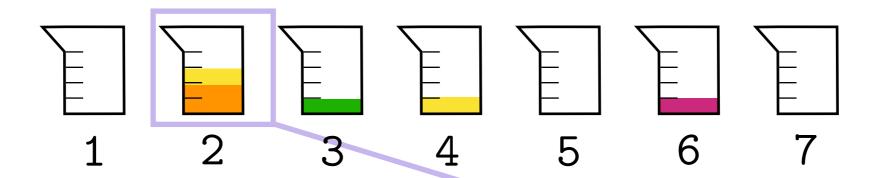


Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second



Then, drain 1 unit from it

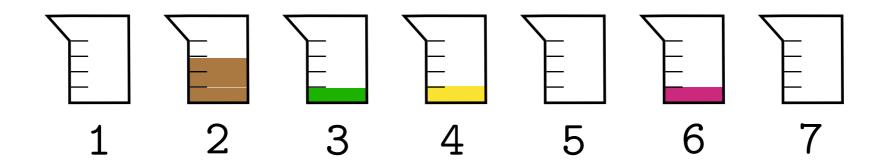


Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second



Then, drain 1 unit from it

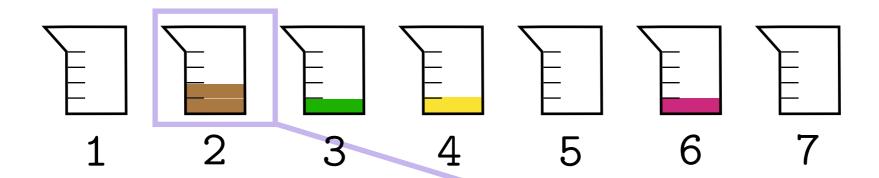


Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it

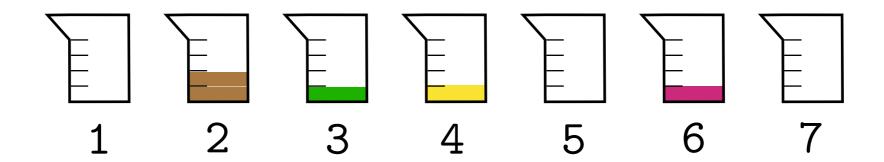


Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it



Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it





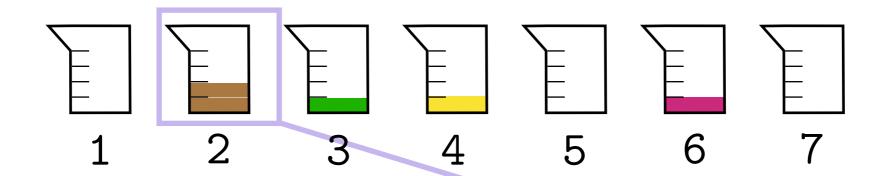
Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it





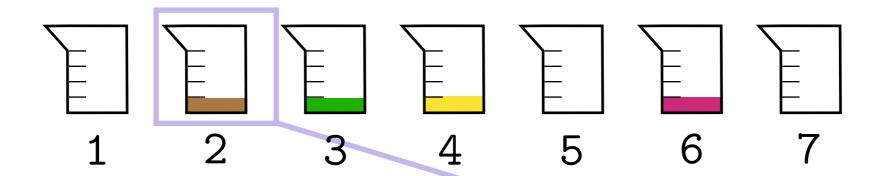
Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it





Empty out the leftmost beaker of purple chemical

Then, add the contents of the first beaker to the second

Mix it

Then, drain 1 unit from it





1. Mix the second beaker

2. Pour it into the third beaker

pour(prevArg1(1), index(allObjects, 3))



1. Mix the second beaker

2. Pour it into the third beaker

```
pour(prevArg1(1), index(allObjects, 3))
```



1. Mix the second beaker

```
mix(index(allObjects, 2))
```

2. Pour it into the third beaker

```
pour(prevArg1(1), index(allObjects, 3))
```



1. Mix the second beaker

Suhr and Artzi 2018

- Sequential understanding: attention over previous instructions
- Grounding in environment: attention over environments, updated as actions are executed
- Talk: Wednesday 18 July, 11:20 (Session 7A)

Outline

- Case study: natural language interface to a database
 - A. Previous instructions
 - B. Previous interpretations of instructions
- 2. Situated interaction context
- 3. Open questions

Open Questions

- Learning by interacting: system questions and user clarifications as a learning signal
- Robustness: training to recover from errors
- Reproducibility and evaluation: working without static corpora
- Latent decisions: learning to distinguish between meaning derived from current utterance vs. interaction history

 Existing work: make use of dialogues to learn context-independent semantic parsers

System | How can I help you?

User I would like to fly from Atlanta, Georgia to London, England on September 24th in the early evening. I would like to return on October 1st departing from London in the late morning.

System | How can I help you?

User I would like to fly from Atlanta, Georgia to London, England on September 24th in the early evening. I would like to return on October 1st departing from London in the late morning.

System Leaving what city?

System failed to recognize city constraint

System | How can I help you?

User I would like to fly from Atlanta, Georgia to London, England on September 24th in the early evening. I would like to return on October 1st departing from London in the late morning.

System Leaving what city?

User Atlanta, Georgia

Can learn from user's response

- Existing work: make use of dialogues to learn context-independent semantic parsers
- Limitations: simple domains; only in contextindependent setting

Robustness

- During training, models should learn how users react to its mistakes
- Learn to recover from mistakes and avoid error propagation
- Existing work: simulated users
- Can collect unlimited number of interactions, but restricted to synthetic language

Reproducibility and Evaluation

- Difficult to compare: No fixed test set when evaluating with live interactions
- Human judgment of understanding is subjective
- Utterance-level evaluation: annotation effort and expertise required
- Interaction-level evaluation: goal completion

Latent Decisions

- Context-independent: how do output tokens align with input utterance?
- Context-dependent: tokens can also align to interaction history
- Large space of latent decisions: is each output token derived from interaction history, or current utterance?
- Existing work: heuristic approaches that place assumptions on how context is used

Summary

- Case study: natural language interface to a database
 - A. Previous instructions
 - B. Previous interpretations of instructions
- 2. Situated interaction context
- 3. Open questions